

Assignment-2

REPORT

Introduction

Here, I tried to run the [Wikipedia-EN-20120601_ARTICLES.tar.gz](#) input-file. But, unfortunately, my laptop has only 8GB RAM, so it wouldn't work out. So, I used [Wikipedia-50-ARTICLES.tar](#) for running the code.

I created the project using Maven, and in the pom.xml file, I added these lines of codes:

```
<properties>
  <maven.compiler.source>8</maven.compiler.source>
  <maven.compiler.target>8</maven.compiler.target>
</properties>

<dependencies>
  <dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-common</artifactId>
    <version>3.4.0</version>
    <scope>system</scope>
    <systemPath>/opt/hadoop-3.4.0/share/hadoop/common/hadoop-common-3.4.0.jar</systemPath>
  </dependency>

  <dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-mapreduce-client-core</artifactId>
    <version>3.4.0</version>
    <scope>system</scope>
    <systemPath>/opt/hadoop-3.4.0/share/hadoop/mapreduce/hadoop-mapreduce-client-core-3.4.0.jar</systemPath>
  </dependency>
</dependencies>
```

Then, copy the relevant java file to the source folder, and run:

```
mvn clean install
```

Now, we have the jar file for the java file. Now, in my system, I had a separate user called "hadoop", which had hadoop in it. Now, in that user, I have to copy the jar file generated. Also, in the HDFS, I have to put the required input file.

Now, run the hadoop command:

```
hadoop@bharli-VivoBook-ASUSLaptop-X515EA-X515EA: $ hadoop jar /home/hadoop/testingCodes/trial-1.0-SNAPSHOT.jar TopWords /testers/Wikipedia-50-ARTICLES.tar /output1/
```

To run hadoop on the jar file, and get the outputs. We can also get some screenshots from localhost:8088.

Co-occurring word matrix generation: Pairs & Stripes and local aggregation

PART-A

Program Logic:

The program aims to identify the top 50 most frequently occurring words from a given input dataset while excluding stop words. It utilizes Hadoop's MapReduce framework to distribute the computation across multiple nodes in a cluster.

Pseudocode Explanation:

Mapper Phase:

Load stop words from the distributed cache.

For each input record:

Tokenize the text into words.

Convert each word to lowercase and check if it is not a stop word.

Emit key-value pairs where the word is the key and the count is set to 1.

Reducer Phase:

Initialize a priority queue to store the top 50 words based on their counts.

For each word received:

Sum up the counts for the same word.

Add the word and its count to the priority queue.

If the queue size exceeds 50, remove the least frequent word.

After processing all words, emit the top 50 words from the priority queue.

Runtime Analysis:

The program utilizes Hadoop's MapReduce framework, which is highly scalable and can handle large datasets efficiently by distributing the workload across multiple nodes.

The Mapper and Reducer tasks are executed in parallel across the cluster, leveraging the distributed computing capabilities of Hadoop.

The runtime performance depends on various factors such as the size of the input dataset, the number of nodes in the cluster, the hardware configuration, and the complexity of the processing logic.

Generally, Hadoop MapReduce jobs are suitable for processing large-scale datasets, and the runtime can scale linearly with the size of the dataset and the cluster's capacity.

Relevant Screenshots:

application_1712340359523_0004	hadoop	top_words	MAPREDUCE	root.default	0	Sun Apr 7 13:40:17 +0550 2024	Sun Apr 7 13:40:17 +0550 2024	Sun Apr 7 13:40:30 +0550 2024	FINISHED	SUCCEEDED	N/A
--------------------------------	--------	-----------	-----------	--------------	---	-------------------------------------	-------------------------------------	-------------------------------------	----------	-----------	-----

Total Resource Preempted:		<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:		0
Total Number of AM Containers Preempted:		0
Resource Preempted from Current Attempt:		<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:		0
Aggregate Resource Allocation:		44015 MB-seconds, 23 vcore-seconds
Aggregate Preempted Resource Allocation:		0 MB-seconds, 0 vcore-seconds

User:	hdoop
Name:	top_words
Application Type:	MAPREDUCE
Application Tags:	
Application Priority:	0 (Higher Integer value indicates higher priority)
YarnApplicationState:	FINISHED
Queue:	root.default
FinalStatus Reported by AM:	SUCCEEDED
Started:	Sun Apr 07 13:40:17 +0530 2024
Launched:	Sun Apr 07 13:40:17 +0530 2024
Finished:	Sun Apr 07 13:40:30 +0530 2024
Elapsed:	12sec
Tracking URL:	History
Log Aggregation Status:	DISABLED
Application Timeout (Remaining Time):	Unlimited
Diagnostics:	
Unmanaged Application:	false
Application Node Label expression:	<Not set>
AM container Node Label expression:	<DEFAULT_PARTITION>

PART-B

Program Logic:

The program aims to construct a co-occurring word matrix based on the input dataset, considering different word distances (d). It utilizes Hadoop's MapReduce framework to distribute the computation across multiple nodes in a cluster.

Pseudocode Explanation:

Mapper Phase:

Tokenize each input record into words.

For each word in the record:

Determine the word distance (d) from the configuration.

Generate pairs of words within the specified distance (excluding the word itself).

Emit key-value pairs where the key is a combination of the word and its co-occurring word, and the value is 1.

Reducer Phase:

Sum up the counts for each word pair.

Emit the word pair and its total count.

Main Function:

Iterate over different word distances specified in the 'distances' array.

Set the word distance in the configuration.

Configure and run the MapReduce job for each word distance.

Measure the runtime for each word distance and print the results.

Runtime Analysis:

The program runs multiple MapReduce jobs, each with a different word distance (d), specified in the 'distances' array.

The runtime performance depends on factors such as the size of the input dataset, the number of nodes in the Hadoop cluster, the complexity of the processing logic, and the specified word distances.

As the word distance increases, the number of word pairs considered for co-occurrence also increases, potentially leading to longer execution times.

The program utilizes Hadoop's distributed processing capabilities to handle large-scale datasets efficiently, with runtime scaling linearly with the dataset size and cluster capacity.

Relevant Screenshots:

application_1712340359523_0008	hadoop	co-occurrence-stripe-d-4	MAPREDUCE	root.default	0	Sun Apr 7 18:47:33 +0550 2024	Sun Apr 7 18:47:39 +0550 2024	Sun Apr 7 18:47:53 +0550 2024	FINISHED	SUCCEEDED	N/A
application_1712340359523_0007	hadoop	co-occurrence-stripe-d-3	MAPREDUCE	root.default	0	Sun Apr 7 18:47:14 +0550 2024	Sun Apr 7 18:47:19 +0550 2024	Sun Apr 7 18:47:32 +0550 2024	FINISHED	SUCCEEDED	N/A
application_1712340359523_0006	hadoop	co-occurrence-stripe-d-2	MAPREDUCE	root.default	0	Sun Apr 7 18:46:54 +0550 2024	Sun Apr 7 18:47:00 +0550 2024	Sun Apr 7 18:47:12 +0550 2024	FINISHED	SUCCEEDED	N/A
application_1712340359523_0005	hadoop	co-occurrence-stripe-d-1	MAPREDUCE	root.default	0	Sun Apr 7 18:46:39 +0550 2024	Sun Apr 7 18:46:40 +0550 2024	Sun Apr 7 18:46:52 +0550 2024	FINISHED	SUCCEEDED	N/A

PART-C

Program Logic:

The program aims to construct a co-occurrence stripe matrix based on the input dataset, considering different word distances (d). It utilizes Hadoop's MapReduce framework to distribute the computation across multiple nodes in a cluster.

Pseudocode Explanation:

Mapper Phase:

Tokenize each input record into words.

For each word in the record:

Determine the word distance (d) from the configuration.

Generate a stripe (map) of co-occurring words within the specified distance (excluding the word itself).

Emit key-value pairs where the key is the word, and the value is the stripe represented as a string.

Reducer Phase:

Merge the stripes for each word, aggregating the counts for co-occurring words.

Emit the word and its consolidated stripe.

Main Function:

Iterate over different word distances specified in the 'distances' array.

Set the word distance in the configuration.

Configure and run the MapReduce job for each word distance.

Measure the runtime for each word distance and print the results.

Runtime Analysis:

The program runs multiple MapReduce jobs, each with a different word distance (d), specified in the 'distances' array.

Similar to the previous example, the runtime performance depends on factors such as the size of the input dataset, cluster resources, and specified word distances.

As the word distance increases, the number of co-occurring words considered in the stripe also increases, potentially leading to longer execution times.

The program utilizes Hadoop's distributed processing capabilities to efficiently handle large-scale datasets, with runtime scaling linearly with the dataset size and cluster capacity.

Relevant Screenshots:

```
Runtime for d = 1: 14483 milliseconds
```

```
File System Counters
  FILE: Number of bytes read=1841365
  FILE: Number of bytes written=4299799
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=768120
  HDFS: Number of bytes written=1783870
  HDFS: Number of read operations=8
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=1905
  Total time spent by all reduces in occupied slots (ms)=1584
  Total time spent by all map tasks (ms)=1905
  Total time spent by all reduce tasks (ms)=1584
  Total vcore-milliseconds taken by all map tasks=1905
  Total vcore-milliseconds taken by all reduce tasks=1584
  Total megabyte-milliseconds taken by all map tasks=1950720
  Total megabyte-milliseconds taken by all reduce tasks=1622016
Map-Reduce Framework
  Map input records=1
  Map output records=112366
  Map output bytes=3072797
  Map output materialized bytes=1841365
  Input split bytes=120
  Combine input records=112366
  Combine output records=26679
  Reduce input groups=26679
  Reduce shuffle bytes=1841365
  Reduce input records=26679
  Reduce output records=26679
  Spilled Records=53358
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=67
  CPU time spent (ms)=3120
  Physical memory (bytes) snapshot=648187904
  Virtual memory (bytes) snapshot=5133131776
  Total committed heap usage (bytes)=537919488
  Peak Map Physical memory (bytes)=390619136
  Peak Map Virtual memory (bytes)=2563297280
  Peak Reduce Physical memory (bytes)=257568768
  Peak Reduce Virtual memory (bytes)=2569834496
```

Runtime for d = 2: 19368 milliseconds

File System Counters

FILE: Number of bytes read=3306555
FILE: Number of bytes written=7230179
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=3243199
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2053
Total time spent by all reduces in occupied slots (ms)=1739
Total time spent by all map tasks (ms)=2053
Total time spent by all reduce tasks (ms)=1739
Total vcore-milliseconds taken by all map tasks=2053
Total vcore-milliseconds taken by all reduce tasks=1739
Total megabyte-milliseconds taken by all map tasks=2102272
Total megabyte-milliseconds taken by all reduce tasks=1780736

Map-Reduce Framework

Map input records=1
Map output records=112366
Map output bytes=5217153
Map output materialized bytes=3306555
Input split bytes=120
Combine input records=112366
Combine output records=26679
Reduce input groups=26679
Reduce shuffle bytes=3306555
Reduce input records=26679
Reduce output records=26679
Spilled Records=53358
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=70
CPU time spent (ms)=3690
Physical memory (bytes) snapshot=672706560
Virtual memory (bytes) snapshot=5131210752
Total committed heap usage (bytes)=588251136
Peak Map Physical memory (bytes)=392941568
Peak Map Virtual memory (bytes)=2565009408
Peak Reduce Physical memory (bytes)=279764992
Peak Reduce Virtual memory (bytes)=2566201344

Runtime for d = 3: 19598 milliseconds

File System Counters

FILE: Number of bytes read=4702419
FILE: Number of bytes written=10021907
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=4633113
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2172
Total time spent by all reduces in occupied slots (ms)=1748
Total time spent by all map tasks (ms)=2172
Total time spent by all reduce tasks (ms)=1748
Total vcore-milliseconds taken by all map tasks=2172
Total vcore-milliseconds taken by all reduce tasks=1748
Total megabyte-milliseconds taken by all map tasks=2224128
Total megabyte-milliseconds taken by all reduce tasks=1789952

Map-Reduce Framework

Map input records=1
Map output records=112366
Map output bytes=7334709
Map output materialized bytes=4702419
Input split bytes=120
Combine input records=112366
Combine output records=26679
Reduce input groups=26679
Reduce shuffle bytes=4702419
Reduce input records=26679
Reduce output records=26679
Spilled Records=53358
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=86
CPU time spent (ms)=3380
Physical memory (bytes) snapshot=696954880
Virtual memory (bytes) snapshot=5129605120
Total committed heap usage (bytes)=631242752
Peak Map Physical memory (bytes)=424828928
Peak Map Virtual memory (bytes)=2562203648
Peak Reduce Physical memory (bytes)=272125952
Peak Reduce Virtual memory (bytes)=2567401472

Runtime for d = 4: 19431 milliseconds

File System Counters

FILE: Number of bytes read=6022984
FILE: Number of bytes written=12663037
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=5946904
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2068
Total time spent by all reduces in occupied slots (ms)=1841
Total time spent by all map tasks (ms)=2068
Total time spent by all reduce tasks (ms)=1841
Total vcore-milliseconds taken by all map tasks=2068
Total vcore-milliseconds taken by all reduce tasks=1841
Total megabyte-milliseconds taken by all map tasks=2117632
Total megabyte-milliseconds taken by all reduce tasks=1885184

Map-Reduce Framework

Map input records=1
Map output records=112366
Map output bytes=9400477
Map output materialized bytes=6022984
Input split bytes=120
Combine input records=112366
Combine output records=26679
Reduce input groups=26679
Reduce shuffle bytes=6022984
Reduce input records=26679
Reduce output records=26679
Spilled Records=53358
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=86
CPU time spent (ms)=3850
Physical memory (bytes) snapshot=715206656
Virtual memory (bytes) snapshot=5128654848
Total committed heap usage (bytes)=637009920
Peak Map Physical memory (bytes)=440426496
Peak Map Virtual memory (bytes)=2561523712
Peak Reduce Physical memory (bytes)=274780160
Peak Reduce Virtual memory (bytes)=2567131136

application_1712340359523_0030	hadoop	co-occurrence-stripe-d-4	MAPREDUCE	root.default	0	Mon Apr 8 07:11:11 +0550 2024	Mon Apr 8 07:11:16 +0550 2024	Mon Apr 8 07:11:28 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0029	hadoop	co-occurrence-stripe-d-3	MAPREDUCE	root.default	0	Mon Apr 8 07:10:52 +0550 2024	Mon Apr 8 07:10:57 +0550 2024	Mon Apr 8 07:11:08 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0028	hadoop	co-occurrence-stripe-d-2	MAPREDUCE	root.default	0	Mon Apr 8 07:10:31 +0550 2024	Mon Apr 8 07:10:37 +0550 2024	Mon Apr 8 07:10:49 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0027	hadoop	co-occurrence-stripe-d-1	MAPREDUCE	root.default	0	Mon Apr 8 07:10:18 +0550 2024	Mon Apr 8 07:10:18 +0550 2024	Mon Apr 8 07:10:30 +0550 2024	FINISHED	SUCCEEDED

PART-D

Program Logic:

CoOccurrenceStripe:

Mapper Logic:

Tokenizes input text into individual words.

Constructs a stripe (a map) for each word, where keys are co-occurring words and values are their counts.

Emits each word along with its corresponding stripe.

Reducer Logic:

Receives word-to-stripe mappings from mappers.

Merges the stripes for each word to generate the final co-occurrence counts.

Emits each word along with its merged stripe.

CoOccurrenceMatrix:

Mapper Logic:

Tokenizes input text into individual words.

Constructs pairs of words within the given distance (d) and emits them with a count of 1.

Reducer Logic:

Receives pairs of words and their counts from mappers.

Aggregates the counts for each pair to generate the final co-occurrence counts.

Emits each pair along with its aggregated count.

Pseudocode Explanations:

CoOccurrenceStripe:

Mapper:

for each word in the input text:

 construct a stripe for the word

 emit (word, stripe)

Reducer:

for each (word, stripe) received:

 merge the stripes for the word

emit (word, merged stripe)

CoOccurrenceMatrix:

Mapper:

for each word in the input text:

 for each co-occurring word within distance d:

 emit (word, co-occurring word) with count 1

Reducer:

for each (word, co-occurring word) received:

 sum up the counts for the pair

emit (word, co-occurring word) with aggregated count

Runtime Analysis:

CoOccurrenceStripe:

Mapper Runtime: $O(n)$, where n is the number of words in the input text.

Reducer Runtime: $O(n)$, where n is the number of distinct words emitted by mappers.

CoOccurrenceMatrix:

Mapper Runtime: $O(n)$, where n is the number of words in the input text.

Reducer Runtime: $O(n^2)$, where n is the number of distinct word pairs emitted by mappers.

Relevant Screenshots:

CoOccurrenceMatrix:

```
Runtime for d = 1: 15110 milliseconds
```

```
File System Counters
  FILE: Number of bytes read=3183188
  FILE: Number of bytes written=6983483
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=768120
  HDFS: Number of bytes written=2575855
  HDFS: Number of read operations=8
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=1994
  Total time spent by all reduces in occupied slots (ms)=1816
  Total time spent by all map tasks (ms)=1994
  Total time spent by all reduce tasks (ms)=1816
  Total vcore-milliseconds taken by all map tasks=1994
  Total vcore-milliseconds taken by all reduce tasks=1816
  Total megabyte-milliseconds taken by all map tasks=2041856
  Total megabyte-milliseconds taken by all reduce tasks=1859584
Map-Reduce Framework
  Map input records=1
  Map output records=224730
  Map output bytes=3948806
  Map output materialized bytes=3183188
  Input split bytes=120
  Combine input records=224730
  Combine output records=151872
  Reduce input groups=151872
  Reduce shuffle bytes=3183188
  Reduce input records=151872
  Reduce output records=151872
  Spilled Records=303744
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=65
  CPU time spent (ms)=2830
  Physical memory (bytes) snapshot=606982144
  Virtual memory (bytes) snapshot=5125009408
  Total committed heap usage (bytes)=507510784
  Peak Map Physical memory (bytes)=374165504
  Peak Map Virtual memory (bytes)=2559729664
  Peak Reduce Physical memory (bytes)=232816640
  Peak Reduce Virtual memory (bytes)=2565279744
```

Runtime for d = 2: 18993 milliseconds

File System Counters

FILE: Number of bytes read=6348046
FILE: Number of bytes written=13313199
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=5150003
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2315
Total time spent by all reduces in occupied slots (ms)=1910
Total time spent by all map tasks (ms)=2315
Total time spent by all reduce tasks (ms)=1910
Total vcore-milliseconds taken by all map tasks=2315
Total vcore-milliseconds taken by all reduce tasks=1910
Total megabyte-milliseconds taken by all map tasks=2370560
Total megabyte-milliseconds taken by all reduce tasks=1955840

Map-Reduce Framework

Map input records=1
Map output records=449458
Map output bytes=7897182
Map output materialized bytes=6348046
Input split bytes=120
Combine input records=449458
Combine output records=299615
Reduce input groups=299615
Reduce shuffle bytes=6348046
Reduce input records=299615
Reduce output records=299615
Spilled Records=599230
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=79
CPU time spent (ms)=3390
Physical memory (bytes) snapshot=659746816
Virtual memory (bytes) snapshot=5128093696
Total committed heap usage (bytes)=596639744
Peak Map Physical memory (bytes)=407576576
Peak Map Virtual memory (bytes)=2560405504
Peak Reduce Physical memory (bytes)=252170240
Peak Reduce Virtual memory (bytes)=2567688192

Runtime for d = 3: 19366 milliseconds

File System Counters

FILE: Number of bytes read=9375936
FILE: Number of bytes written=19368979
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=7620158
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2612
Total time spent by all reduces in occupied slots (ms)=2204
Total time spent by all map tasks (ms)=2612
Total time spent by all reduce tasks (ms)=2204
Total vcore-milliseconds taken by all map tasks=2612
Total vcore-milliseconds taken by all reduce tasks=2204
Total megabyte-milliseconds taken by all map tasks=2674688
Total megabyte-milliseconds taken by all reduce tasks=2256896

Map-Reduce Framework

Map input records=1
Map output records=674184
Map output bytes=11844860
Map output materialized bytes=9375936
Input split bytes=120
Combine input records=674184
Combine output records=439058
Reduce input groups=439058
Reduce shuffle bytes=9375936
Reduce input records=439058
Reduce output records=439058
Spilled Records=878116
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=81
CPU time spent (ms)=4030
Physical memory (bytes) snapshot=702513152
Virtual memory (bytes) snapshot=5130289152
Total committed heap usage (bytes)=596639744
Peak Map Physical memory (bytes)=450252800
Peak Map Virtual memory (bytes)=2559926272
Peak Reduce Physical memory (bytes)=252260352
Peak Reduce Virtual memory (bytes)=2570362880

Runtime for d = 4: 22552 milliseconds

File System Counters

FILE: Number of bytes read=12251066
FILE: Number of bytes written=25119239
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=9967043
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=4200
Total time spent by all reduces in occupied slots (ms)=3317
Total time spent by all map tasks (ms)=4200
Total time spent by all reduce tasks (ms)=3317
Total vcore-milliseconds taken by all map tasks=4200
Total vcore-milliseconds taken by all reduce tasks=3317
Total megabyte-milliseconds taken by all map tasks=4300800
Total megabyte-milliseconds taken by all reduce tasks=3396608

Map-Reduce Framework

Map input records=1
Map output records=898908
Map output bytes=15792304
Map output materialized bytes=12251066
Input split bytes=120
Combine input records=898908
Combine output records=571164
Reduce input groups=571164
Reduce shuffle bytes=12251066
Reduce input records=571164
Reduce output records=571164
Spilled Records=1142328
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=122
CPU time spent (ms)=7010
Physical memory (bytes) snapshot=710627328
Virtual memory (bytes) snapshot=5128368128
Total committed heap usage (bytes)=638058496
Peak Map Physical memory (bytes)=439087104
Peak Map Virtual memory (bytes)=2560798720
Peak Reduce Physical memory (bytes)=271540224
Peak Reduce Virtual memory (bytes)=2567569408

CoOccurrenceStripe:

Runtime for d = 1: 13078 milliseconds

```
File System Counters
  FILE: Number of bytes read=1021573
  FILE: Number of bytes written=2660239
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=768120
  HDFS: Number of bytes written=966650
  HDFS: Number of read operations=8
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=1800
  Total time spent by all reduces in occupied slots (ms)=1550
  Total time spent by all map tasks (ms)=1800
  Total time spent by all reduce tasks (ms)=1550
  Total vcore-milliseconds taken by all map tasks=1800
  Total vcore-milliseconds taken by all reduce tasks=1550
  Total megabyte-milliseconds taken by all map tasks=1843200
  Total megabyte-milliseconds taken by all reduce tasks=1587200
Map-Reduce Framework
  Map input records=1
  Map output records=112366
  Map output bytes=2736428
  Map output materialized bytes=1021573
  Input split bytes=120
  Combine input records=112366
  Combine output records=26679
  Reduce input groups=26679
  Reduce shuffle bytes=1021573
  Reduce input records=26679
  Reduce output records=26679
  Spilled Records=53358
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=67
  CPU time spent (ms)=2800
  Physical memory (bytes) snapshot=629624832
  Virtual memory (bytes) snapshot=5129814016
  Total committed heap usage (bytes)=508559360
  Peak Map Physical memory (bytes)=384655360
  Peak Map Virtual memory (bytes)=2564984832
  Peak Reduce Physical memory (bytes)=244969472
  Peak Reduce Virtual memory (bytes)=2564829184
```

Runtime for d = 2: 18788 milliseconds

File System Counters

FILE: Number of bytes read=2557535
FILE: Number of bytes written=5732163
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=2497211
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=1995
Total time spent by all reduces in occupied slots (ms)=1669
Total time spent by all map tasks (ms)=1995
Total time spent by all reduce tasks (ms)=1669
Total vcore-milliseconds taken by all map tasks=1995
Total vcore-milliseconds taken by all reduce tasks=1669
Total megabyte-milliseconds taken by all map tasks=2042880
Total megabyte-milliseconds taken by all reduce tasks=1709056

Map-Reduce Framework

Map input records=1
Map output records=112366
Map output bytes=4663514
Map output materialized bytes=2557535
Input split bytes=120
Combine input records=112366
Combine output records=26679
Reduce input groups=26679
Reduce shuffle bytes=2557535
Reduce input records=26679
Reduce output records=26679
Spilled Records=53358
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=83
CPU time spent (ms)=3350
Physical memory (bytes) snapshot=678043648
Virtual memory (bytes) snapshot=5131993088
Total committed heap usage (bytes)=566755328
Peak Map Physical memory (bytes)=426008576
Peak Map Virtual memory (bytes)=2565177344
Peak Reduce Physical memory (bytes)=252035072
Peak Reduce Virtual memory (bytes)=2566815744

Runtime for d = 3: 20308 milliseconds

File System Counters

FILE: Number of bytes read=3953343
FILE: Number of bytes written=8523779
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=3887180
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2273
Total time spent by all reduces in occupied slots (ms)=1828
Total time spent by all map tasks (ms)=2273
Total time spent by all reduce tasks (ms)=1828
Total vcore-milliseconds taken by all map tasks=2273
Total vcore-milliseconds taken by all reduce tasks=1828
Total megabyte-milliseconds taken by all map tasks=2327552
Total megabyte-milliseconds taken by all reduce tasks=1871872

Map-Reduce Framework

Map input records=1
Map output records=112366
Map output bytes=6566824
Map output materialized bytes=3953343
Input split bytes=120
Combine input records=112366
Combine output records=26679
Reduce input groups=26679
Reduce shuffle bytes=3953343
Reduce input records=26679
Reduce output records=26679
Spilled Records=53358
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=82
CPU time spent (ms)=3540
Physical memory (bytes) snapshot=680169472
Virtual memory (bytes) snapshot=5128638464
Total committed heap usage (bytes)=595591168
Peak Map Physical memory (bytes)=401051648
Peak Map Virtual memory (bytes)=2559119360
Peak Reduce Physical memory (bytes)=279117824
Peak Reduce Virtual memory (bytes)=2569519104

Runtime for d = 4: 20769 milliseconds

File System Counters

FILE: Number of bytes read=5272821
FILE: Number of bytes written=11162735
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=768120
HDFS: Number of bytes written=5201021
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=3136
Total time spent by all reduces in occupied slots (ms)=2431
Total time spent by all map tasks (ms)=3136
Total time spent by all reduce tasks (ms)=2431
Total vcore-milliseconds taken by all map tasks=3136
Total vcore-milliseconds taken by all reduce tasks=2431
Total megabyte-milliseconds taken by all map tasks=3211264
Total megabyte-milliseconds taken by all reduce tasks=2489344

Map-Reduce Framework

Map input records=1
Map output records=112366
Map output bytes=8425191
Map output materialized bytes=5272821
Input split bytes=120
Combine input records=112366
Combine output records=26679
Reduce input groups=26679
Reduce shuffle bytes=5272821
Reduce input records=26679
Reduce output records=26679
Spilled Records=53358
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=128
CPU time spent (ms)=6380
Physical memory (bytes) snapshot=681979904
Virtual memory (bytes) snapshot=5128122368
Total committed heap usage (bytes)=597688320
Peak Map Physical memory (bytes)=405123072
Peak Map Virtual memory (bytes)=2559590400
Peak Reduce Physical memory (bytes)=276856832
Peak Reduce Virtual memory (bytes)=2568531968

application_1712340359523_0035	hadoop	co-occurrence-matrix-d-4	MAPREDUCE	root.default	0	Mon Apr 8 08:52:39 +0550 2024	Mon Apr 8 08:52:45 +0550 2024	Mon Apr 8 08:53:01 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0034	hadoop	co-occurrence-matrix-d-3	MAPREDUCE	root.default	0	Mon Apr 8 08:52:20 +0550 2024	Mon Apr 8 08:52:25 +0550 2024	Mon Apr 8 08:52:38 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0033	hadoop	co-occurrence-matrix-d-2	MAPREDUCE	root.default	0	Mon Apr 8 08:52:01 +0550 2024	Mon Apr 8 08:52:06 +0550 2024	Mon Apr 8 08:52:18 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0032	hadoop	co-occurrence-matrix-d-1	MAPREDUCE	root.default	0	Mon Apr 8 08:51:46 +0550 2024	Mon Apr 8 08:51:47 +0550 2024	Mon Apr 8 08:51:59 +0550 2024	FINISHED	SUCCEEDED

application_1712340359523_0043	hadoop	co-occurrence-stripe-d-4	MAPREDUCE	root.default	0	Mon Apr 8 09:05:47 +0550 2024	Mon Apr 8 09:05:52 +0550 2024	Mon Apr 8 09:06:06 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0042	hadoop	co-occurrence-stripe-d-3	MAPREDUCE	root.default	0	Mon Apr 8 09:05:27 +0550 2024	Mon Apr 8 09:05:33 +0550 2024	Mon Apr 8 09:05:45 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0041	hadoop	co-occurrence-stripe-d-2	MAPREDUCE	root.default	0	Mon Apr 8 09:05:09 +0550 2024	Mon Apr 8 09:05:13 +0550 2024	Mon Apr 8 09:05:25 +0550 2024	FINISHED	SUCCEEDED
application_1712340359523_0040	hadoop	co-occurrence-stripe-d-1	MAPREDUCE	root.default	0	Mon Apr 8 09:04:56 +0550 2024	Mon Apr 8 09:04:56 +0550 2024	Mon Apr 8 09:05:06 +0550 2024	FINISHED	SUCCEEDED

Indexing Documents via Hadoop

PART-A

Program Logic:

The program aims to calculate the document frequency of terms in a corpus using Hadoop's MapReduce framework. It preprocesses the input text by removing stop words and performing stemming before emitting each term with its document ID and frequency of occurrence.

Pseudocode Explanation:

Mapper Phase:

Load stop words from the "stopwords.txt" file into a set.

For each input record:

Tokenize the text into words.

Lowercase each word and remove non-alphanumeric characters.

Skip stop words and perform stemming using the Porter Stemmer.

Emit each term with its document ID (key) and value 1.

Reducer Phase:

Aggregate the frequency counts for each term across different documents.

Emit each term along with its total frequency in the corpus.

Main Function:

Configure the MapReduce job with the necessary classes, jar file, and input/output formats.

Set the mapper, combiner (optional), and reducer classes.

Specify the output key and value classes.

Set the input and output paths.

Run the job and exit with status 0 if successful, else exit with status 1.

Runtime Analysis:

The program's runtime performance depends on factors such as the size of the input corpus, the number of terms, and the efficiency of the stemming algorithm.

Preprocessing steps like stop word removal and stemming may introduce overhead, but they contribute to more accurate document frequency calculations by eliminating noise from the text.

The MapReduce framework enables parallel processing of input data across multiple nodes in a cluster, which helps handle large-scale corpora efficiently.

Relevant Screenshots:

Got a stemmer error, so couldn't get the final output...

User:	hadoop
Name:	document frequency
Application Type:	MAPREDUCE
Application Tags:	
Application Priority:	0 (Higher Integer value indicates higher priority)
YarnApplicationState:	FINISHED
Queue:	root.default
FinalStatus Reported by AM:	FAILED
Started:	Mon Apr 08 06:48:48 +0530 2024
Launched:	Mon Apr 08 06:48:48 +0530 2024
Finished:	Mon Apr 08 06:49:04 +0530 2024
Elapsed:	16sec
Tracking URL:	History
Log Aggregation Status:	DISABLED
Application Timeout (Remaining Time):	Unlimited
Diagnostics:	Task failed task_1712340359523_0020_m_000000 Job failed as tasks failed. failedMaps:1 failedReduces:0 killedMaps:0 killedReduces: 0
Unmanaged Application:	false
Application Node Label expression:	<Not set>
AM container Node Label expression:	<DEFAULT_PARTITION>

Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	54454 MB-seconds, 28 vcore-seconds
Aggregate Preempted Resource Allocation:	0 MB-seconds, 0 vcore-seconds

PART-B

Program Logic:

The program calculates the TF-IDF (Term Frequency-Inverse Document Frequency) scores for terms in a corpus using the "stripes" approach. It first computes the term frequency (TF) for each term within each document and then calculates the IDF component using the total number of documents in the corpus. Finally, it combines these values to compute the TF-IDF score for each term.

Pseudocode Explanation:

Mapper Phase:

Parse each input record, which consists of a document ID followed by term-count pairs.

Emit each document ID along with its corresponding term-count pair.

Reducer Phase:

For each document ID, aggregate the term-count pairs and calculate the total number of terms in the document.

Compute the TF-IDF score for each term in the document using the formula:

$$\text{TF-IDF} = \text{term frequency} / \text{total terms}$$

Emit each document ID along with its corresponding term and TF-IDF score.

Main Function:

Configure the MapReduce job with the necessary classes, jar file, and input/output formats.

Set the mapper and reducer classes.

Specify the output key and value classes.

Set the input and output paths.

Run the job and exit with status 0 if successful, else exit with status 1.

Runtime Analysis:

The program's runtime performance depends on factors such as the size of the corpus, the number of unique terms, and the efficiency of the TF-IDF computation.

The MapReduce framework enables parallel processing of input data across multiple nodes in a cluster, which helps handle large-scale corpora efficiently.

The complexity of the TF-IDF calculation is linear with respect to the number of unique terms in each document and the total number of documents in the corpus.

Relevant Screenshots:

application_1712340359523_0022	hadoop	TF-IDF Stripes	MAPREDUCE	root.default	0	Mon Apr 8 06:55:58 +0550 2024	Mon Apr 8 06:55:58 +0550 2024	Mon Apr 8 06:56:08 +0550 2024	FINISHED	SUCCEEDED
--------------------------------	--------	-------------------	-----------	--------------	---	-------------------------------------	-------------------------------------	-------------------------------------	----------	-----------

```
File System Counters
  FILE: Number of bytes read=6
  FILE: Number of bytes written=616353
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=768120
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=8
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=1441
  Total time spent by all reduces in occupied slots (ms)=1406
  Total time spent by all map tasks (ms)=1441
  Total time spent by all reduce tasks (ms)=1406
  Total vcore-milliseconds taken by all map tasks=1441
  Total vcore-milliseconds taken by all reduce tasks=1406
  Total megabyte-milliseconds taken by all map tasks=1475584
  Total megabyte-milliseconds taken by all reduce tasks=1439744
```